Q4 · Reflection and Next Steps

Mira · Interview Kickstart, Applied Agentic AI for PMs/TPMs · Ashutosh Lonkar · November 2026

What the traces showed

- **Grounding works because code does the maths.** Across 24 test runs (v1 and v2), every count, task ID, risk ID and date matched the files. Re-runs changed wording, never facts.
- **Traces caught what reading answers would miss.** An 80× inflated count exposed a Download node re-reading files per item; the routing field on each trace confirmed 12 / 12 requests reached the intended agent.
- **Evaluation found a real gap.** v1 status reports lacked upcoming milestones, a core requirement. v2 fixed it; the regression run proved no other test changed.
- **What remains is wording, not facts:** a counts-only question answered with a full report (T10), filler like "significant progress" (T12), risk ranking that varies between runs (T3 vs T7), and milestone answers that do not link tasks to the milestone they feed (T11).
- **Cost is negligible** ($0.0005 per request, about $0.83 a month for 10 PMs); latency (~10 s) is dominated by two LLM calls and four file downloads.

Improvement plan

- **v3 prompts:** fix the six logged defects, each with a before/after regression run compared by tag in Langfuse.
- **Grounding validator:** a Code node after each agent checks every ID and date in the answer against the files and scores the trace automatically.
- **Rank in code, write in LLM:** compute risk priority from task-board evidence so rankings are stable.
- **Human in the loop:** PM approval (Approve/Reject) before the weekly email reaches stakeholders.
- **Speed and reliability:** own Google OAuth app and a daily file cache to remove quota errors and cut latency.

From evaluation to fine-tuning

Every trace already carries input, output and two human scores. Answers scored below 4 are corrected by a PM and saved to a **Langfuse dataset** as input / corrected-output pairs. Once about **50 corrected examples** exist for the Risk Assessor (the agent with the most judgment), they become a fine-tuning set for gpt-4o-mini. The fine-tuned model is accepted only if it passes all 12 regression tests, keeps grounding at 100% and raises average usefulness. Prompt fixes come first; fine-tuning is for consistent style and judgment, not for facts, which stay in code.

Privacy considerations

- **What is sensitive:** staff names and workload on the task board, client details and risks, and stakeholder emails.
- **Minimise:** each agent already receives only its own facts; traces will store IDs and redact staff names and client identifiers.
- **Control where data goes:** an LLM contract with no training on customer data; self-hosted Langfuse inside Nexora; 30-day trace retention; Drive folders restricted to the project team.
- **Approve before sharing:** nothing reaches a client without PM review; new data sources need Data Protection sign-off.

Rollout to 10 PMs/TPMs

- **Weeks 1–4, pilot:** 2 PM champions use Mira for weekly status and stakeholder updates; time saved and usefulness tracked. Go / no-go: 100% accuracy, usefulness ≥ 4, report time < 10 min.
- **Weeks 5–10, expand to 10:** 30-minute onboarding, a playbook of example requests, a Drive folder per project, PM approval on every outgoing email, weekly office hours.
- **Operate:** weekly review of low-scoring traces; monthly regression run and cost report; success = 8 of 10 PMs using Mira weekly and 10–20 PM-hours saved per week.